



CANCER
RESEARCH
UK

CAMBRIDGE
INSTITUTE

CRUK CI cluster introduction (III of III)

Some advanced topics



UNIVERSITY OF
CAMBRIDGE

Reference Genomes

- Path to reference data:
`/scratchb/bioinformatics/reference_data/reference_genomes/`
- Path to assembly:
.../organism/assembly/
- What we maintain:
 - Genome sequence (fasta)
 - Alignment indices: BWA, TopHat, Bowtie (1,2)
 - Annotations:
 - GTF format gene model
 - RefFlat format gene model
 - Signal artifact list (if available)

Working with Lustre

1. Revisit architecture
2. Stripes
3. Avoiding I/O Bottlenecks
4. Using System Cache

Lustre: Quick Review

Lustre is a massively parallel distributed file system

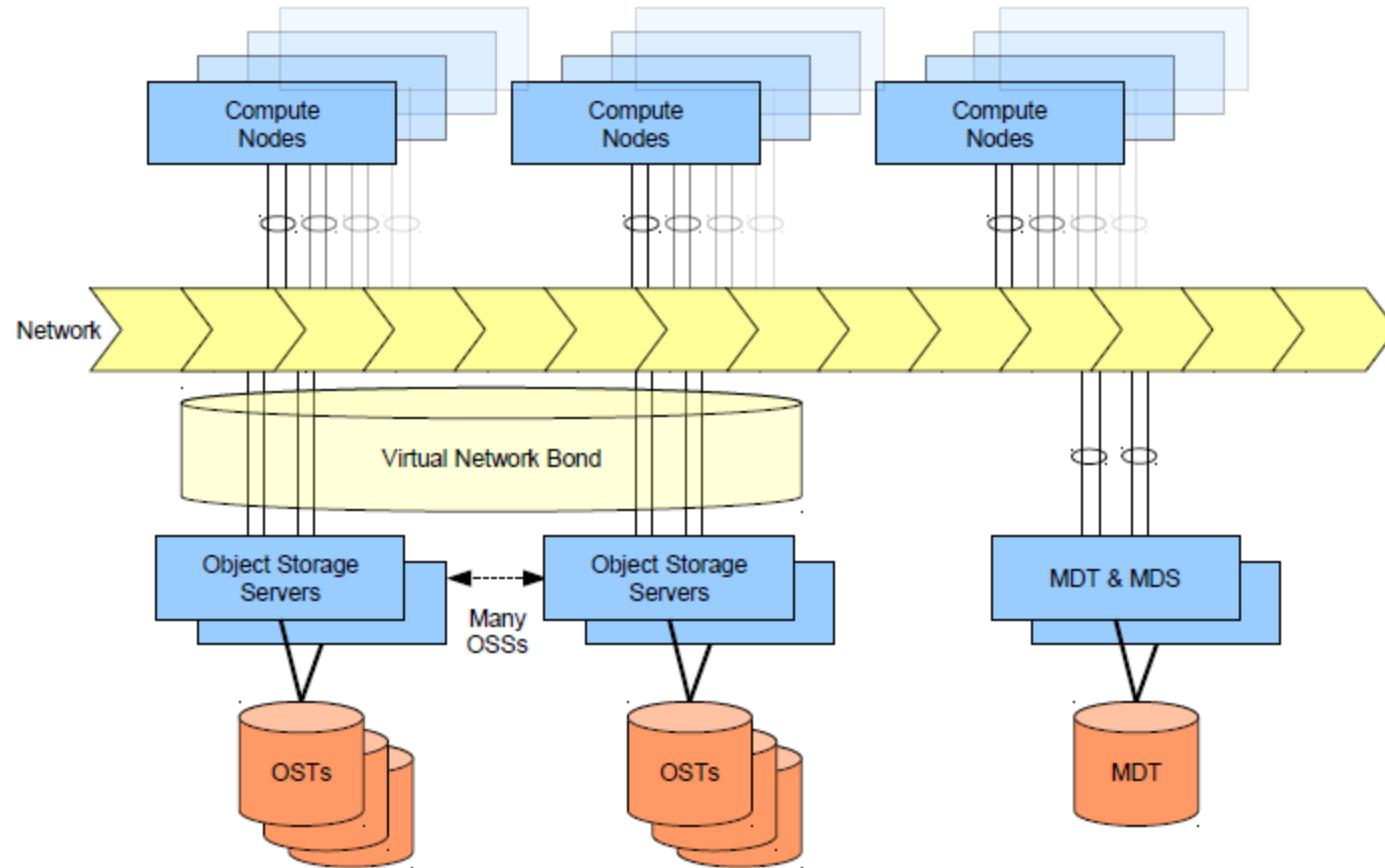
- Deployed in 7 out of 10 most powerful supercomputers
- POSIX compliant

Lustre design paradigm concepts

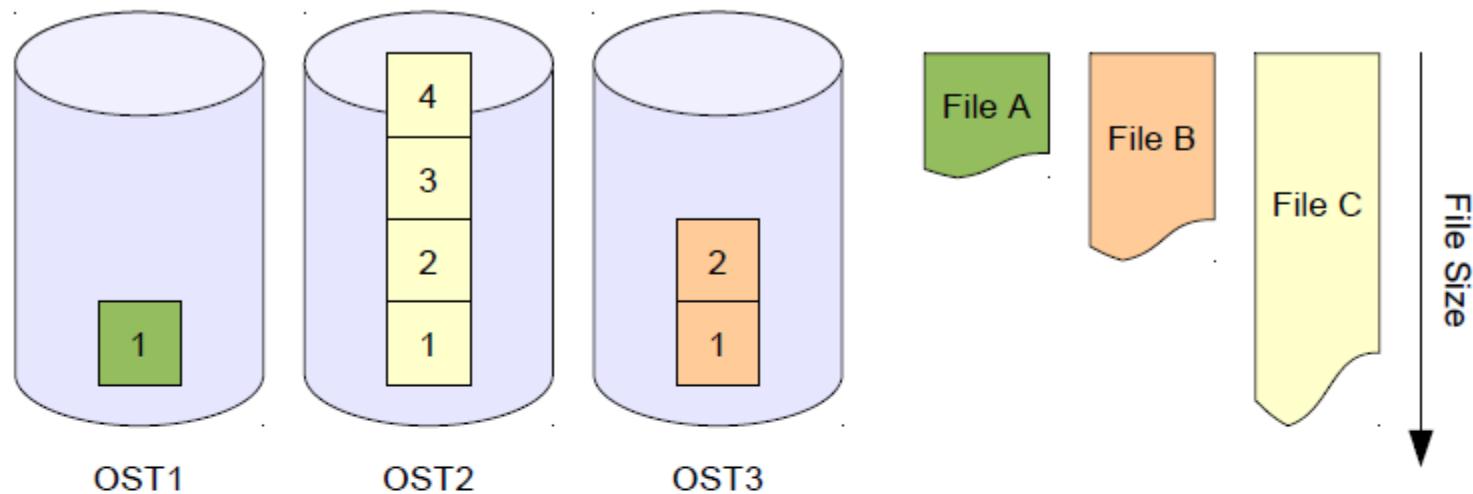
- Separation of file meta-data and storage allocation
- Scalable data serving through parallel data striping
- Aggregates network bandwidth
- Distributed operation

'scratch' storage we (deliberately) don't back it up

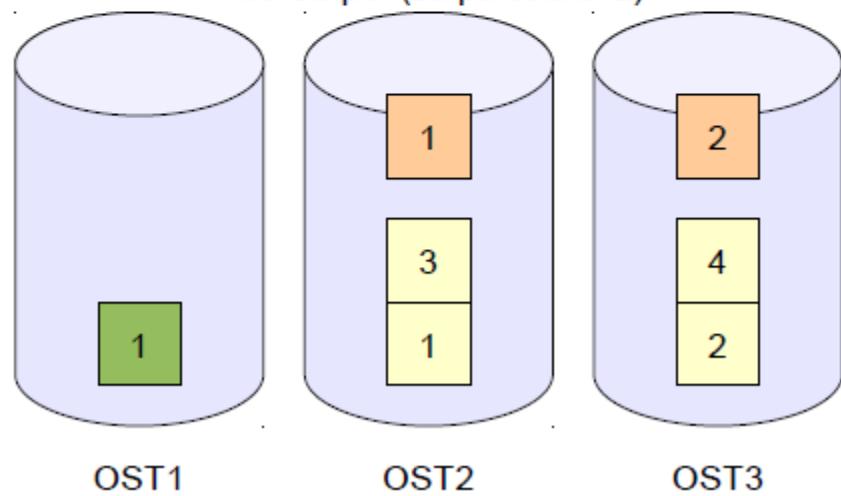
Lustre Architecture ...



Single Stripe (stripe count=1)



Two Stripe (stripe count=2)



...and File Striping

File Striping Large Files

Striping allows file size to exceed single OST size

Performance benefit

- Aggregates I/O bandwidth to single large file
- In general, more strips improves performance
- Small overhead associated with open/closing striped files

Many jobs reading single file

- For example blastdb and maq reference data
- /lustre/reference_data/genomes

Many jobs reading & writing multiple large files

- Requires benchmarking

Many jobs writing to single file

- High bandwidth but requires careful coding (can be disastrous)

Set Stripe Information

Set per file or directory

Default is not to stripe

Only newly created files will be stripes

- Use cp (not mv) to migrate existing files

```
clust1-headnode ~ $ lfs setstripe <file|dir> --size <stripe_size>
--count <count> --index <index>
```

Where,

size = stripe size specified in k, m or g (0 default 1MB)

count = OST stripe count (0 defaults 4 OST and -l over all OSTs)

index = OST index of first stripe (-l indicating default)

Read Stripe Information

Inspect file and directory stripe information
with **lfs getstripe**

```
uk-cri-lcst01 ~ $ lfs getstripe -d /lustre/reference_data/genomes
stripe_count: 65535 stripe_size: 0 stripe_offset: 0
uk-cri-lcst01 ~ $ lfs getstripe -d /lustre/reference_data
stripe_count: 1 stripe_size: 1048576 stripe_offset: 0
uk-cri-lcst01 ~ $ lfs getstripe zma.3.ebwt
/lustre/reference_data/genomes/Zea_mays/zma.3.ebwt
lmm_stripe_count: 16
lmm_stripe_size: 1048576
lmm_stripe_offset: 11
obdidx objid objid group
11 623006 0x9819e 0
0 8504376 0x81c438 0
...
12 622252 0x97eac 0
8 607894 0x94696 0
```

Using System Cache

Disk access is slow (no escape from this!)

- Memory access measured in a few nanoseconds
- Disc access measured in 10s of milliseconds

Linux uses free memory as cache

- Memory reclaimed as least used files expunged

“Pre-warming” cache

- Can increase I/O performance

```
uk-cri-lcst01 ~ $ cat largefile > /dev/null  
uk-cri-lcst01 ~ $ grep searchString largefile
```

Avoiding Cache misses

Say you want to compare 3 sequences against 3 large databases:

“Out of the box” example

- sequence 1 vs database 1 **No cache: disk read required**
- sequence 1 vs database 2 **Cache miss: disk read required**
- sequence 1 vs database 3 **Cache miss: disk read required**
- sequence 2 vs database 1 **No cache: data expunged from cache**
- sequence 2 vs database 2 **Cache miss: disk read required**

Re-ordering to avoid cache misses

- sequence 1 vs database 1 **No cache: disk read required**
- sequence 2 vs database 1 **Cache hit: data in cache**
- sequence 3 vs database 1 **Cache hit: data in cache**
- sequence 1 vs database 2 **No cache: disk read required**
- sequence 2 vs database 2 **Cache hit: data in cache**

Avoiding Bad Performance

- Interactive use
 - Stattting files can be slow (common with shared file systems)
 - Avoid directly editing small files on lustre (keep to /home)
 - Turn off “color ls” (stat required for each file/directory)

Random seeks

- Small random I/O extremely slow on lustre
- Avoid, as much as possible, running databases on lustre
e.g. mysql, sqlite, Berkeley DB etc

Limit number of files in directory

- 10,000s files in single directory bad (avoid, if possible)
- Use lfs setstripe to confine all files to single OST – obviously for small files only!

```
uk-cri-lcst01 ~ $ lfs setstripe --count 1 directory
```

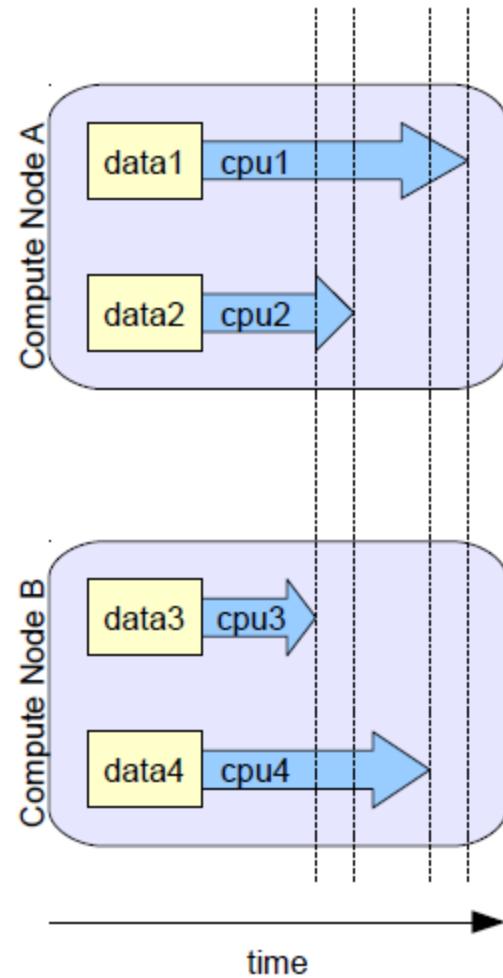
Simple Parallel

(‘Embarassing’ or ‘Trivial’ in the computing science literature)

Solving many similar and independent tasks

- Analysis split into tasks
- Task assigned to one cpu
- No inter-task communication
- More throughput by running more tasks
- Task runtime varies

90% of bioinformatics codes fall into this model



Shared Memory

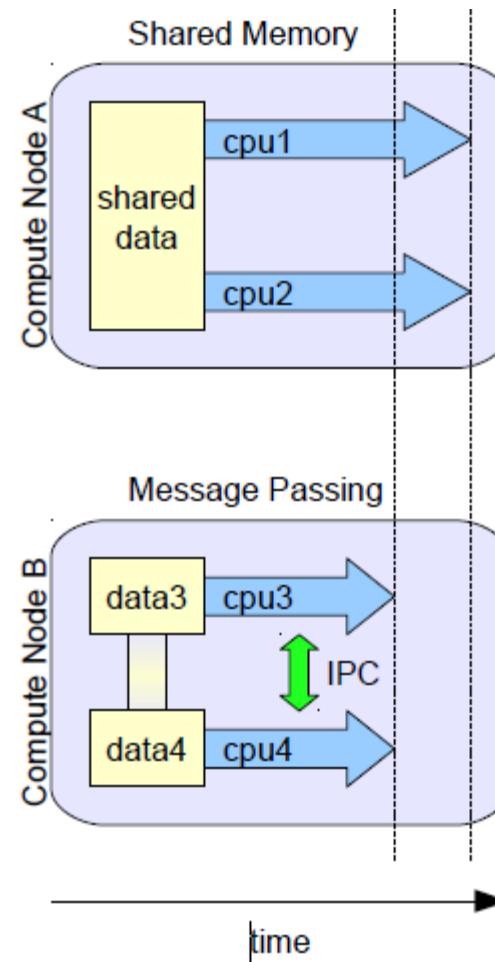
Shared tasks and memory

- Tasks assigned to cpus or cores
- Inter-task communication via shared memory
- Runtime decreases by adding more threads

Message Passing - local

Multiple processes communicate using O/S level systems. Code must be specifically written to exploit parallelism

- OpenMP/OpenMPI/etc



Message passing over network

Single task split across many compute nodes

- inter-machine communication (IMC)
- through MPI/OpenMP libraries again

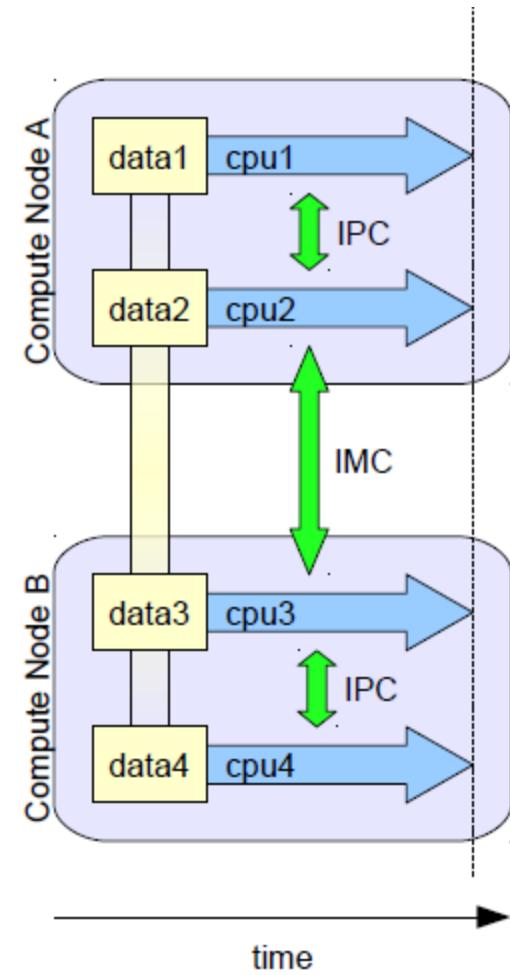
Hybrid models

Single task split across many compute nodes

- Mix SMP, local MP, network MP

Can be tricky to predict performance.

Your code may get quite complex...



Parallel Workloads

What happens when you increase your dataset size?

- $O(n)$, $O(n \log n)$, $O(n^2)$, $O(n^9)$
- Runtime & memory increase (O) with problem size (n)

Calculations are carried out in parallel

- Operating on principle that large calculations can often be divided into N smaller tasks
- These tasks are solved concurrently
- Time reduces to a function of $O(n^x)/N$

Parallel overheads

- communication, concurrency, parallel I/O introduce new overheads
- $O(N)$, $O(N \log N)$, $O(N^2)$

Amdahl's Law

- Speed up of parallel application is ultimately limited by the fixed runtime of any sequential sections
- There's no magic bullet for this, you may have to change your algorithm, or mix parallel and single node sections.



CANCER
RESEARCH
UK

CAMBRIDGE
INSTITUTE

Acknowledgements

Peter MacCallum
Marc O'Brien
Mark Dunning
Ann Pajon



UNIVERSITY OF
CAMBRIDGE